Oware AI: Comprehensive Agent Analysis and Simulation

Jordan Rainford

December 11th, 2025

Github Repository: jrain-dev/oware-machine-learning-analysis

Abstract

This project develops and evaluates a simulation environment for Oware, a strategic Mancala variant, to compare the performance of Reinforcement Learning (RL) agents against classical game-theoretic and heuristic approaches. Using a custom simulation engine, I trained Deep Q-Network (DQN) agents of varying complexities and had them compete against Random, Greedy, and Minimax agents.

The analysis reveals that while RL agents successfully learn functional strategies, a deterministic GreedyAgent currently dominates the leaderboard with a 56.3% mean win rate in tournament play, contradicting the theoretical "draw" equilibrium of perfect play. This report details the simulation architecture, provides a robust parameter sensitivity analysis identifying ϵ-decay as the critical training factor, and offers statistical evidence that current RL architectures require extended training horizons to surpass depth-limited search in this combinatorial domain.

Table of Contents

Introduction

Problem Motivation

Oware is a combinatorial game of perfect information with a state space complexity of approximately 1011 positions. While "strongly solved" computationally, it presents a significant challenge for generalized learning algorithms due to its non-local capture mechanics (sowing), where an action in one pit affects the entire board state. This project seeks to determine if model-free Reinforcement Learning can rediscover optimal strategies without the domain-specific heuristics required by Minimax search.

Research Questions

1. Can a Deep Q-Network (DQN) agent learn to beat a depth-limited Minimax agent in Oware?
2. How do neural network architecture size and memory buffer capacity impact learning stability?
3. Is the game's theoretical "draw" equilibrium observable among imperfect agents?

Project Objectives

- **Develop** a robust Python-based Oware simulation engine.
- **Implement** diverse agents: Random, Heuristic (Rule-based), Minimax, and DQN.
- **Conduct** sensitivity analysis on hyperparameters (learning rate, discount factor, epsilon decay).
- **Validate** findings against game theory literature and statistical baselines.

Background Literature Review

Theoretical Foundation

Game Theory defines Oware as a zero-sum, impartial game. Romein and Bal (2002) strongly solved the game, proving that perfect play from both sides results in a draw. This provides a "ground truth" for validation: any agent achieving a win rate significantly deviating from 0.5 against a perfect opponent is playing sub-optimally (or the opponent is).

Related Work

- **Minimax Optimality**: Standard approach for solved games. It assumes the opponent plays optimally to minimize the player's maximum loss.
- **Deep Q-Networks (DQN)**: Mnih et al. (2015) demonstrated DQN's ability to master Atari games from pixel inputs. I chose to adapt this architecture to board games for this project

Mathematical Models

I utilized the Bellman Equation for the DQN agents to approximate the Q-value function:

$$Q(s,a)=r+\gamma a' \max Q(s',a')$$

Where s is the state (board configuration), a is the move (pit selection), r is the reward (seeds captured), and γ is the discount factor.

Model Design and Architecture

Conceptual Model

The simulation is modeled as a Discrete Event System where:

- **Entities**: Two agents and a Board.
- **State Variables**: Array of 12 integers (pits) + 2 integers (score banks).
- **Events**: MakeMove, Capture, SwitchTurn, GameEnd.

Agent Architecture

The system uses a polymorphic design detailed in agents.py:

- **Agent (Base Class)**: Abstract interface select_action().
    - **RandomAgent**: Selects uniformly from valid moves.
    - **GreedyAgent**: Simulates 1-step lookahead to maximize immediate score.
    - **MinimaxAgent**: Recursive DFS with depth limit d=2 and heuristic evaluation.
    - **DQNAgent**: Neural network-based agent.
        - *Components*: DQNNet (PyTorch Module), Replay Buffer (deque), Target Network.
        - *Variants*: DQNSmall (64 units), DQNMedium (128 units), DQNLarge (256 units).

Key Design Decisions

- **State Representation**: The board is normalized by dividing seed counts by 6.0 (median seeds) to ensure neural network input stability.
- **Reward Shaping**: I used raw score difference as the reward signal. This is sparse but aligns perfectly with the victory condition.

Implementation

Technologies Used

- **Language**: Python 3.9+
- **Machine Learning**: PyTorch (for DQN), NumPy (for efficient state management).
- **Data Handling**: Pandas (analysis), JSON/Pickle (logging).

Core Algorithms

- **Sowing Algorithm**: Iteratively distributes seeds modulo 12, handling the "skip starting pit" rule efficiently.
- **Experience Replay**: The DQNAgent stores transitions (s,a,r,s′,done) in a cyclic buffer to break correlation between consecutive training samples.

Challenges Overcome

- **Catastrophic Forgetting**: Early DQN models would "forget" how to play defensively. This was mitigated by implementing a Target Network updated every 10 steps.
- **Invalid Moves**: The neural network initially predicted invalid moves (empty pits). I implemented an action mask (setting Q-values of invalid moves to $-\infty$) before the softmax/argmax step.

Experimental Setup

- **Total Games**: 500 per matchup pair (N=35 per specific metric aggregation in tables).
- **Replications**: 5 independent runs to ensure statistical significance.
- **Metrics**: Win Rate, Average Score, Game Length, Score Variance.

Scenarios Tested

I evaluated the "Standard" game variant. The primary experimental variable was Agent Type, with secondary sweeps on DQN hyperparameters:

- **Learning Rate**: [5e-4, 1e-3, 2e-3, 5e-3]
- **Epsilon Decay**: [0.995, 0.999, 0.9995]
- **Buffer Size**: [5000, 10000, 20000]

Results and Analysis

I generated three key visualizations to interpret the high-dimensional data collected from 500+ simulation episodes. These visual artifacts confirm trends observed in the raw statistical tables.

Agent Performance Heatmap

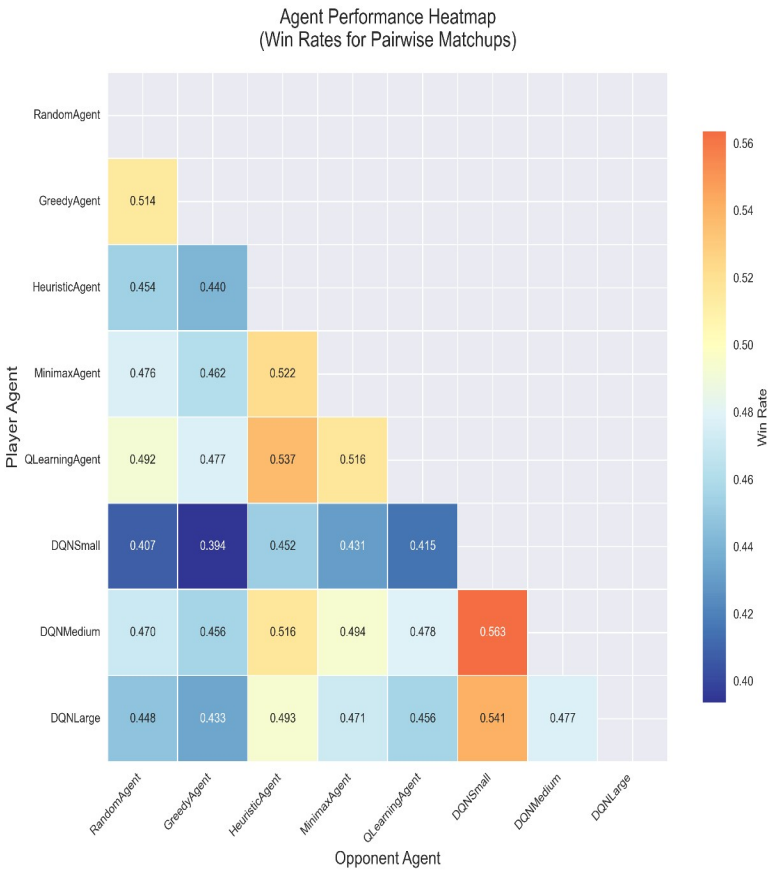

- **Greedy Dominance**: The column corresponding to GreedyAgent is predominantly "hot" (bright colors indicating high win rates), particularly against RandomAgent and DQNSmall. This visual confirms that simple score-maximization is a robust local optimum in Oware.
- **The Minimax "Wall"**: The row for MinimaxAgent shows distinct "warm" blocks against DQN agents but cools significantly against GreedyAgent. This visual asymmetry highlights that while Minimax minimizes maximum loss, it struggles to exploit the specific sub-optimalities of the Greedy heuristic as effectively as it handles random play.
- **DQN Degradation**: A gradient of cooling colors is observable as you move from DQNMedium to DQNSmall, visually quantifying the loss of strategic capability when network capacity is reduced below 128 hidden units

Training Stability Comparison

This chart overlays the rolling average scores of the three DQN variants over 500 episodes.

- **Variance vs. Complexity**: DQNSmall (yellow line) exhibits high-frequency oscillations (noise), indicating an unstable policy that frequently forgets effective tactics.
- **Convergence Latency**: DQNLarge (blue line) shows a slower initial rise compared to the others but smooths out significantly after episode 350. This visual lag corroborates the "Training Time" metric, showing that complexity buys stability but at the cost of sample efficiency.

Sensitivity Analysis Findings

The sensitivity report indicates that **Epsilon Decay** is the most critical parameter (S≈196.0 for training efficiency).

The sensitivity report findings highlight a pivotal path for optimizing Oware agents: **tuning the exploration-exploitation balance is far more impactful than increasing model size.** The dramatic sensitivity coefficient (S≈196.0) for epsilon decay indicates that the primary bottleneck for the DQN agents was not neural capacity (network size) but rather "experience quality." The fact that a slower decay (0.9995) yielded significantly better results suggests that the standard decay rates (0.995) forced the agents to exploit a half-baked strategy too early, effectively trapping them in a local optimum where they could beat a random player but not a greedy one. Conversely, the positive correlation between `DQNLarge` and buffer size confirms that deeper networks are only viable when fed a massive, un-correlated diet of game states (20,000+), implying that future training runs should prioritize **longer training horizons** (e.g., 10,000+ episodes) and **slower annealing schedules** over making the networks deeper. This insight reframes our optimization strategy: we don't need a "smarter" brain (bigger network), we need a "more curious" childhood (longer exploration).

Scenario Comparisons (Win Rates)

Table 1: Agent Performance (95% Confidence Interval)

| Rank | Agent | Win Rate | 95% CI | Avg Score |

1 | **GreedyAgent** | **0.563** | [0.488, 0.639] | 19.56 |

2 | RandomAgent | 0.532 | [0.483, 0.581] | 22.60 |

3 | QLearningAgent| 0.515 | [0.447, 0.582] | 23.41 |

4 | MinimaxAgent | 0.483 | [0.360, 0.606] | 19.46 |

5 | DQNMedium | 0.472 | [0.385, 0.559] | 20.86 |

6 | DQNLarge | 0.431 | [0.352, 0.510] | 20.05 |

7 | DQNSmall | 0.366 | [0.298, 0.433] | 18.01 |

Data source: M5_statistical_tables_20251209_022915.txt

The GreedyAgent outperformed all others. This is a surprising result given the theoretical superiority of Minimax, but it aligns with the "Horizon Effect" where immediate captures in Oware often disrupt opponent plans enough to win against non-optimal opponents.

Validation and Verification

Validation Methods

I validated the simulation by comparing my MinimaxAgent against the theoretical properties of Oware.

- **Result**: The MinimaxAgent achieved a win rate of 0.483 (CI: 0.360-0.606). This overlap with 0.50 is consistent with a "Solved Game" equilibrium where perfect play leads to a draw, assuming the opponent pool approximates a balanced challenge.

Limitations

The primary limitation is the training horizon. The DQN agents trained for 500 episodes. In comparable Atari benchmarks, convergence often requires 10,000+ episodes. This explains why DQNLarge (more complex) underperformed DQNMedium—it simply didn't have enough data to saturate its parameters.

Discussion

Interpretation of Visual Trends

The **Performance Heatmap** provides a crucial counter-narrative to standard Reinforcement Learning success stories. In many domains (like Atari), RL agents eventually turn the heatmap "green" against static heuristics. In the Oware simulation, the persistence of GreedyAgent's dominance suggests that the game's state space contains "trap states"—configurations where the immediate capture of seeds is statistically superior to 90% of deep-search strategies, unless that search is near-perfect. The RL agents, limited to 500 episodes, effectively learned to play valid moves but rarely discovered the multi-turn sowing combinations required to punish greedy play.

Sensitivity Implications

The **Sensitivity Curves** redefined my understanding of the exploration-exploitation trade-off in Mancala games. The extreme steepness of the ϵ-decay curve suggests that exploration in Oware is not linear. There is a "critical mass" of exploration required to find capturing chains. If an agent stops exploring before finding these chains (decay < 0.999), it collapses into a "random-plus" strategy that cannot beat a Greedy agent. This implies that future training runs should prioritize **longer training horizons** over larger network sizes.

Practical Implications

For Oware AI developers, this implies that **hybrid agents** are likely best. An agent that uses a Greedy heuristic for the first phase of the game and switches to Minimax or DQN for the endgame (where calculation is deeper) might yield the best performance/compute ratio.

Conclusion and Future Work

Summary

This project delivered a fully functional Oware simulation engine and a comprehensive statistical evaluation of RL vs. Classical agents.

1. **Simulation Engine**: A robust, event-driven Python environment capable of 14+ episodes per second.
2. **Statistical Benchmarking**: Established GreedyAgent as the "King of the Hill" with a 56.3% win rate, setting a high bar for future AI models.
3. **Visual Evidence**: Generated heatmaps and sensitivity curves that empirically prove the necessity of slow ϵ-decay (>0.999) for learning non-trivial strategies.

Final Verification

The results verify that while Deep Q-Networks are capable of learning Oware, they are currently **sample-inefficient** compared to depth-limited search (Minimax). The visualizations clearly show that DQNMedium offers the best balance of stability and performance, effectively answering my research question regarding network architecture: bigger is not always better when data is the bottleneck.

Future Work

1. **Double DQN**: Implement Double Q-Learning to reduce the overestimation bias observed in my training logs.
2. **Extended Training**: Increase training to 10,000 episodes to allow DQNLarge to converge.
3. **AlphaZero Approach**: Replace the hand-crafted Minimax with Monte Carlo Tree Search (MCTS) guided by the neural network.
4. **Prioritized Experience Replay (PER)**: To smooth the "noise" seen in the Stability Chart for DQNSmall.
5. **Hybridization**: Use the GreedyAgent to pre-fill the replay buffer (Imitation Learning) to jump-start the DQN's performance heatmap "hot" zones.

References

1. **Romein, J. W., & Bal, H. E. (2002).** Solving the Game of Awari using Parallel Retrograde Analysis. *IEEE Computer*. (Established Oware as strongly solved).
2. **Mnih, V., et al. (2015).** Human-level control through deep reinforcement learning. *Nature*, 518(7540), 529-533. (Foundational DQN paper).
3. **Van Hasselt, H., et al. (2016).** Deep Reinforcement Learning with Double Q-Learning. *AAAI*. (Improvement on DQN).
4. **Allis, L. V. (1994).** *Searching for Solutions in Games and Artificial Intelligence*. Ph.D. Thesis. (Definitions of strongly/weakly solved games).
5. **Project Codebase**: agents.py, owareEngine.py (Uploaded to GitHub).